

Automated Quality Control of Facade Elements

Tim - Fasada na Oko

1 Intro

KFK (Krov, Fasada, Konstrukcija) is a European leader in aluminum and glass facade systems. Each facade element must pass a quality control (QC) step before installation. Today, this process is manual: a worker visually inspects each element and compares it against specifications. This approach is slow, error-prone, and expensive—missed defects are often discovered only on-site, where repair or replacement costs are extremely high.

This project was developed during **NeoData Hackathon 2025**, organized by **Commminus**, with the goal of exploring an automated, scalable computer vision solution for defect detection on facade elements.

2 Approach / Solution

2.1 Dataset Exploration

We started by thoroughly exploring the provided dataset and available 3D models of facade elements. The first step was manual inspection of images to understand:

- Typical appearance of correct (positive) elements
- Common defect patterns in defective (negative) elements
- Variability in lighting, viewpoints, and background clutter

This step was crucial since the dataset contained **no predefined defect annotations**.

2.2 Early Experiments and Key Insight

We initially explored **SAM3** (Segment Anything Model v3) to directly localize defects. While SAM3 showed promising qualitative results, its performance was inconsistent without additional structural context.

Defect detection is significantly easier if the facade element type is known first.

The dataset contains **three distinct facade element types**, each with different geometry and defect characteristics.

2.3 Macro Idea (Final System Design)

We defined the following high-level pipeline:

1. **Detect and classify the facade element** in the image
2. **Apply a specialized defect-detection strategy** tailored to that specific facade element type

This allowed us to trade a single, generic defect detector for **multiple precise, domain-specific methods**.

2.4 Dataset Limitations and Synthetic Data Exploration

The main challenge was the **small size of the real dataset**. We explored generating a synthetic dataset using the provided **3D models** by rendering realistic images with controlled defects. Although conceptually strong, this approach proved too time-consuming to implement robustly within the 24-hour hackathon.

As a practical compromise, we:

- Applied image rotations and simple augmentations
- Balanced classes to stabilize training

2.5 Facade Element Detection (YOLO)

We trained a **YOLO-based object detection model** to classify the facade element type.

- Primary metric: **accuracy**
- Rationale: downstream defect detection depends entirely on correct element classification

This model acts as the entry point to the pipeline.

2.6 Defect Detection (SAM3 + Heuristics)

After classification, we applied **element-specific defect detection strategies**, combining:

- Carefully designed **SAM3 prompts**
- Geometry- and context-aware post-processing

Here, **precision** was the dominant metric: false positives are costly in QC and directly penalized in evaluation.

2.7 Explored but Dropped: 3D–2D Alignment

We also investigated aligning the ideal **3D model** with the **2D image** to detect deviations as defects. While promising for long-term industrial deployment, accurate 3D–2D alignment exceeded the scope of a 24-hour hackathon and was therefore dropped.

2.8 Productization

To make the solution usable, we implemented a **simple Streamlit UI** that:

- Accepts a single image or a batch of images
- Runs the full pipeline
- Generates a clear QC report per image

2.9 Why Now?

KFK is actively expanding its **Data Science and automation efforts**. This solution demonstrates a scalable path toward:

- Reducing manual QC costs
- Improving consistency and reliability
- Laying foundations for future data-driven quality assurance systems

The approach is modular and can evolve with more data and tighter integration of 3D models.

2.10 Product Demo

The final solution is wrapped in a lightweight **Streamlit application** designed for fast and intuitive QC inspection.

The app supports:

- **Single-image analysis** — upload one image and receive an immediate QC decision and defect report
- **Batch analysis** — upload a .zip archive of images and automatically generate a consolidated report for all elements

For each image, the application:

- Predicts the facade element type
- Runs the corresponding defect-detection pipeline
- Generates a structured QC report (PASS / FAIL with detected defects)
- **Visually highlights selected defects** directly on the image for easier interpretation

The demo includes screenshots of example usage for both single-image and batch workflows.

2.11 Demo Screenshots

Facade Defect Detector

Single image → immediate report. ZIP upload → CSV report for all images inside.

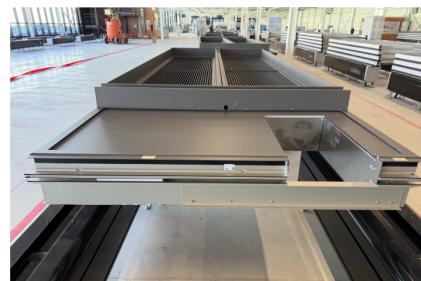
[Single image](#) [Batch \(ZIP\)](#)

Upload an image

Drag and drop file here
Limit 200MB per file • JPG, JPEG, PNG, BMP, WEBP

Browse files

 IMG_5628.jpg 5.7MB



Report

Image: IMG_5628.jpg
Status: FAIL
Defects: screw missing

x

Facade Defect Detector

Single image → immediate report. ZIP upload → CSV report for all images inside.

[Single image](#) [Batch \(ZIP\)](#)

Upload an image

Drag and drop file here
Limit 200MB per file • JPG, JPEG, PNG, BMP, WEBP

Browse files

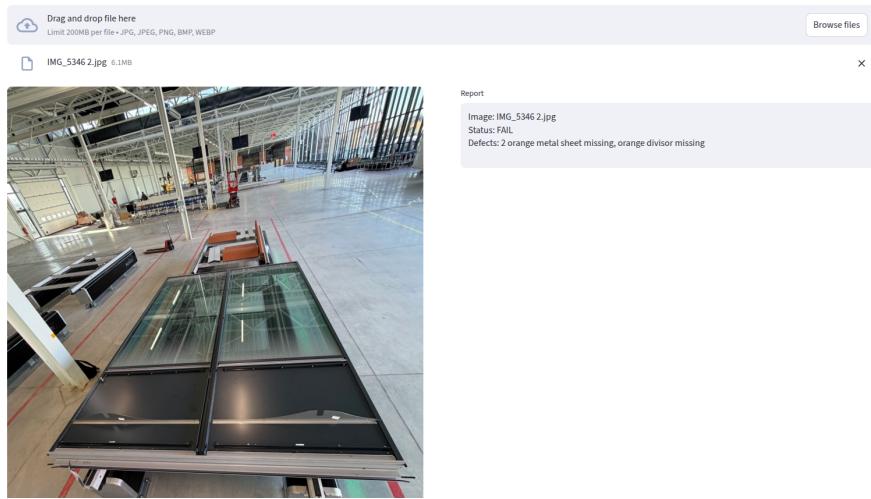
 IMG_9827.jpg 0.9MB



Report

Image: IMG_9827.jpg
Status: FAIL
Defects: broken glass

x



2.12 Long-Term Vision

Although not fully implemented during the hackathon, **3D-model-based synthetic data generation** remains a highly promising direction. With sufficient time, this approach could enable:

- Large, realistic training datasets
- Controlled simulation of rare defects
- Strong generalization across projects and configurations

2.13 Team

- **Marin Meštrović** — Student, PMF (Graduate Program in Computer Science)
- **Damjan Crnković** — Student, FER
- **Luka Bubnjević** — Student, FER